

Zumin Li

SOFTWARE ENGINEER · FULL STACK DEVELOPER

☎ (+86) 13675823946 | ✉ zumin.li.work@gmail.com | 🏠 resume.flyingjack.top | 🌐 FlyingJackLee | 📞 13675823946

INTRODUCTION

Brings 5+ years of software development experience and is currently focused on AI Agent application development, with practical knowledge of agentic patterns including ReAct, Plan-and-Execute and Reflection, as well as core mechanisms such as task planning, state management, tool use, memory and context management. Able to build agents and workflows using LangChain, LangGraph, LangChain4j, MCP and RAG. Possesses a strong foundation in Java backend engineering, distributed microservices and cloud-native systems, alongside full-stack development, architecture design, performance optimization and containerized delivery. Capable of independently delivering the complete engineering lifecycle from solution design and implementation to deployment and operations.

SKILLS

AI Agent Development	LangGraph, LangChain, LangChain4j, ReAct, Plan-and-Execute, Reflection, Human-in-the-loop
LLM Engineering	MCP, RAG, Tool Calling, LLM Gateway, Prompt Engineering, Embedding, pgvector
Backend	Java, Python, Spring Boot, Spring Cloud Alibaba, FastAPI, Kafka, REST API
Cloud Native & Delivery	Kubernetes, Docker, Istio, Jenkins, GitOps, Prometheus
Data & Client	PostgreSQL, MySQL, Redis, Elasticsearch, React, Angular, Flutter

EDUCATION

MSc in Software Engineering - Distinction
UNIVERSITY OF GLASGOW

Glasgow, UK
September 2020 – April 2022

BEng in Communication Engineering - Outstanding Graduate
HANGZHOU DIANZI UNIVERSITY

Hangzhou, China
September 2012 – June 2016

WORK

China Telecom Corporation Limited
FORWARD DEPLOYED ENGINEER

Chongqing, China
March 2026 – Present

- Embedded with enterprise customers to translate ambiguous AI and MaaS requirements into concrete, testable use cases, delivering end-to-end PoCs and demos within 1–2 weeks. Iterated rapidly on customer feedback to move initiatives from exploration to implementation.
- Productized validated PoCs into AI agents embedded in real business workflows, owning workflow abstraction, agent architecture and core implementation. Built task orchestration, tool calling, retrieval and multi-agent collaboration using LangGraph, FastAPI, RAG and MCP.
- Owned end-to-end delivery from prototype to production, including MaaS integration, enterprise system integration, observability, containerization and Kubernetes deployment. Continuously optimized prompts, retrieval and agent strategies based on real-world usage and on-site feedback.

Qilin Communications Co., Ltd.
TECHNICAL LEAD & PARTNER

Guiyang, China
November 2023 – February 2026

- Led full-stack development of an intelligent integrated warehouse and cashier management platform, focusing on Java backend and distributed microservice architecture.
- Optimized microservice architecture and system performance, and delivered iterative upgrades with cloud-native technologies. Managed Kubernetes container orchestration, automated deployment, environment configuration and version release, significantly enhancing system stability, high-concurrency performance, scalability and maintainability.

Huawei Cloud Computing Technologies Co., Ltd
SENIOR SOFTWARE ENGINEER

Chengdu, China
July 2022 – August 2023

- Developed and iterated backend features for commercial PaaS cloud testing products, ensuring high availability and high-performance architecture to support stable commercial delivery.
- Independently led the full-cycle design and development of an in-house testing framework, implementing automated test data injection and API Mock simulation to significantly improve team testing efficiency.
- Provided on-site technical support and end-to-end project risk management for hybrid cloud customers, guaranteeing smooth project deployment and operation.

Zhejiang Science Education Broadcasting Production Co., Ltd.

Hangzhou, China

SOFTWARE ENGINEER

August 2017 – August 2019

- Delivered full-stack development of the enterprise OA approval system and corporate website, covering workflow modelling, backend APIs, frontend interactions and cross-device adaptation to improve configurability and maintainability.

PROJECTS

Expert Engine – Industry Intelligence Reasoning Engine

Release @ August 2026

SYSTEM ARCHITECTURE & CORE DEVELOPMENT

Overview: An AI expert engine for business opportunity intelligence, built with FastAPI, PostgreSQL, pgvector and OpenSearch. It comprises Expert Profile, Knowledge Engine, Organization Engine, Need Engine, Capability Engine, Evidence Engine, Scoring Engine, Reviewer and a knowledge learning loop. The engine automatically identifies opportunities from policies and news, recommends target organizations and departments, matches company capabilities, and provides explainable scores and actionable recommendations.

Responsibilities:

- Designed the overall architecture and core modules, including the Expert Profile framework, hybrid retrieval pipeline (vector + BM25 + SQL) and evidence-driven reasoning workflow to ensure explainability and traceability.
- Developed the Organization Engine and Need Engine, implementing intelligent recommendation of organizations and departments based on responsibility matching, topic mapping and historical cases, with department roles (lead/support) and need maturity classification.

AI Bid Assistant

Release @ July 2026

AI AGENT & FULL-STACK DEVELOPMENT · Github: github.com/FlyingJackLee/bid-agent

Overview: A LangGraph-based bidding assistant covering tender parsing, compliance review, human verification, batch evaluation and bidder comparison, with integrated OCR, regulation-focused RAG and automated reporting.

Responsibilities:

- Designed stateful LangGraph workflows for document parsing, requirement extraction, compliance checking and report generation, with Human-in-the-loop review and checkpoint-based recovery.
- Built a unified tool and security layer for document processing, OCR, data sanitization and PDF generation, preventing unsanitized content from reaching LLM and embedding services.
- Developed regulation-focused RAG with PostgreSQL, pgvector and local embeddings, and productized the agent with FastAPI, React and TypeScript for concurrent batch review, human decisions and progress tracking.

Enterprise LLM Data Privacy Gateway

Release @ June 2026

SYSTEM ARCHITECTURE & CORE DEVELOPMENT

Overview: An LLM security gateway for enterprise AI Agents and MaaS services, using local security processing and cloud inference to detect, tokenize and restore sensitive data within the enterprise boundary.

Responsibilities:

- Designed the hybrid architecture and LLM Gateway for unified access to AI applications, security policies, private models and cloud MaaS services.
- Combined rule engines, named entity recognition and local LLMs for semantic detection, dynamic sanitization and security validation.
- Designed the Token Vault and core services for data replacement, mapping, request proxying, response restoration and access control.

Smart Inventory and Cashier Management Platform

Release @ November 29, 2024

FULL STACK DEVELOPMENT · Github: <https://github.com/flyingjack-cloud>

Overview: An integrated platform for warehouse retail operations, covering inventory, order settlement, cashier verification and mobile workflows. The system evolved from a monolith into a cloud-native microservice architecture based on Spring Cloud Alibaba, Kubernetes and Istio, with AI Agent capabilities for contextual business orchestration.

Responsibilities:

- Delivered full-stack development across domain, API and database design, built the Angular and Flutter clients, and integrated business agents with Langchain4j.
- Led three phases of architectural evolution, introducing distributed caching and the Istio service mesh for routing, traffic control, circuit breaking, canary releases and A/B testing.
- Built a Jenkins and GitOps CI/CD pipeline integrating code checks, image builds, multi-environment deployment and version rollback.

Overview: A PaaS cloud testing product within Huawei Cloud CodeArts DevOps, providing test execution, data management and service support for commercial and hybrid-cloud customers, with an emphasis on high availability, performance and reliable delivery in complex dependency environments.

Responsibilities:

- Developed a new test data injection tool, enabling automated multi-level injection of mock data in complex dependency environments, cutting the team's development cycle by over 30%.
- Implemented backend feature development based on Java/Python and microservice architecture, designed and maintained standardized REST APIs, and optimized business logic for distributed systems to guarantee high availability and performance.

See my website for more: <https://resume.flyingjack.top/>